

MAD Exam 2019

MAD 2018
Department of Computer Science
University of Copenhagen

Exam number: 46

Version 1

Due: January 20th, 23:59

Contents

1. Exercise	3
(a)	3
(b)	3
(c)	4
2. Exercise	5
(a)	5
(b)	6
3. Exercise	6
(a)	6
(b)	8
(c)	9

1. Exercise

(a)

To solve this, we need to realise that a pdf always has the following property: It's integral over the entire function is equal to 1. Since we know that

$$f(x) = 0, \text{ if } 0 > x, 5 < x$$

We can compute the definite integral between 0 and 5, $\int_0^5 f(x)dx$ to find the integral. This means that we have to solve the following equation:

$$\int_0^5 c(5-x)dx = 1$$

First we will extend it to make it easier to work with

$$\int_0^5 5c - xcdx = 1$$

We know that definite integrals are solved as follows.

$$\int_a^b f(x)dx = F(b) - F(a)$$

First we find $F(x)$:

$$\int 5c - xcdx = 5xc - \frac{x^2c}{2}$$

Now we solve $F(5) - F(0)$:

$$(5 * 5c - \frac{5^2c}{2}) - (5 * 0c - \frac{0^2c}{2}) = 1$$

This can be simplified to:

$$12.5c = 1$$

We can then realise that for this to be true, $c = \frac{2}{25}$

(b)

First we will find the *pdf* of V . We know that to find the *pdf*, we have to compute the derivative of the *cdf*. To do this on a piecewise function, we compute the derivative of each piece:

$$f(v) = \frac{d}{dv}F(v) = \begin{cases} \frac{d}{dv}0 & \text{if } v < 0 \\ \frac{d}{dv}v^5 & \text{if } 0 \leq v \leq 1 \\ \frac{d}{dv}1 & \text{if } v > 1 \end{cases} = \begin{cases} 0 & \text{if } v < 0 \\ 5v^4 & \text{if } 0 \leq v \leq 1 \\ 0 & \text{if } v > 1 \end{cases}$$

To find the expectation, $\mathbb{E}V$ we have to use the following formula, since V is a continuous random variable. We know that $\int f(v)dv = F(v)$:

$$\mathbb{E}V = \int vf(v) = \begin{cases} \int 0v dv & \text{if } v < 0 \\ \int 5v^5 dv & \text{if } 0 \leq v \leq 1 \\ \int 0v dv & \text{if } v > 1 \end{cases} = \begin{cases} 0 & \text{if } v < 0 \\ \frac{5v^6}{6} & \text{if } 0 \leq v \leq 1 \\ 0 & \text{if } v > 1 \end{cases}$$

(c)

To prove this, we first have to set an $\epsilon > 0$. We know the following:

$$\mathbb{P}(|Y_n - Y| \geq \epsilon) \rightarrow 0$$

$$\mathbb{P}(|Z_n - Z| \geq \epsilon) \rightarrow 0$$

And we want to prove that:

$$\mathbb{P}(|(Y_n - Y) + (Z_n - Z)| \geq \epsilon) \rightarrow 0$$

We start off by writing the following, which we know is true, since the value of ϵ is arbitrarily small:

$$\mathbb{P}\left(|Y_n - Y| \geq \frac{\epsilon}{2}\right) \rightarrow 0$$

$$\mathbb{P}\left(|Z_n - Z| \geq \frac{\epsilon}{2}\right) \rightarrow 0$$

We can now write the following:

$$\mathbb{P}\left(|Y_n - Y| \geq \frac{\epsilon}{2}\right) + \mathbb{P}\left(|Z_n - Z| \geq \frac{\epsilon}{2}\right) \rightarrow 0$$

We can then write the following, and also see that it still approaches 0:

$$\mathbb{P}\left(|Y_n - Y| \geq \frac{\epsilon}{2}\right) + \mathbb{P}\left(|Z_n - Z| \geq \frac{\epsilon}{2}\right) \geq \mathbb{P}\left(|Z_n - Z| \geq \frac{\epsilon}{2} \vee |Y_n - Y| \geq \frac{\epsilon}{2}\right)$$

$$\mathbb{P}\left(|Z_n - Z| \geq \frac{\epsilon}{2} \vee |Y_n - Y| \geq \frac{\epsilon}{2}\right) \rightarrow 0$$

We can see that this probability *must* be greater than $\mathbb{P}(|(Y_n - Y) + (Z_n - Z)| \geq \epsilon)$, since for this to be true at least one of the events in the previous probability have to happen. Because of this, we can write the following:

$$\mathbb{P}(|(Y_n - Y) + (Z_n - Z)| \geq \epsilon) \leq \mathbb{P}\left(|Z_n - Z| \geq \frac{\epsilon}{2} \vee |Y_n - Y| \geq \frac{\epsilon}{2}\right)$$

We can now say that since

$$\mathbb{P}\left(|Z_n - Z| \geq \frac{\epsilon}{2} \vee |Y_n - Y| \geq \frac{\epsilon}{2}\right) \rightarrow 0$$

and

$$\mathbb{P}(|(Y_n - Y) + (Z_n - Z)| \geq \varepsilon) \leq \mathbb{P}\left(|Z_n - Z| \geq \frac{\varepsilon}{2} \vee |Y_n - Y| \geq \frac{\varepsilon}{2}\right)$$

The following must be true:

$$\mathbb{P}(|(Y_n - Y) + (Z_n - Z)| \geq \varepsilon) \longrightarrow 0$$

This is exactly what we wanted to prove, since:

$$Y_n + Z_n \xrightarrow{\mathbb{P}} Y + Z = \mathbb{P}(|(Y_n - Y) + (Z_n - Z)| \geq \varepsilon) \longrightarrow 0$$

2. Exercise

(a)

First, we will find the likelihood, which is the same as the *joint pdf*, which is defined as thus:

$$f_{\theta}^{joint}(x_1, \dots, x_n) = \prod_{i=1}^n f_{\theta}(x_i)$$

We can then figure out that the likelihood is:

$$f_{\beta}^{joint}(x_1, \dots, x_n) = \begin{cases} \prod_{i=1}^n f_{\beta}(x_i) & \text{if } 0 \leq x \leq \beta \\ 0 & \text{otherwise.} \end{cases} = \begin{cases} \prod_{i=1}^n \frac{2}{\beta^2} \cdot (\beta - x_i) & \text{if } 0 \leq x \leq \beta \\ 0 & \text{otherwise.} \end{cases}$$

We then have to figure out the maximum likelihood estimator. This is done using the following formula:

$$\hat{\theta}_n^{ML}(x_1, \dots, x_n) = \underset{\theta}{argmax} \prod_{i=1}^n f_{\theta}(x_i)$$

We insert the previous formula with the given values for X :

$$\hat{\beta}_n^{ML}(x_1, \dots, x_n) = \underset{\beta}{argmax} \prod_{i=1}^n f_{\beta}(x_i) = \underset{\beta}{argmax} \left(\frac{2}{\beta^2} \cdot (\beta - 3) \cdot \frac{2}{\beta^2} \cdot (\beta - 4) \right)$$

First we rewrite it to make our lives easier:

$$\frac{4(\beta - 4)(\beta - 3)}{\beta^4}$$

Now we have to differentiate it to find the local maxima:

$$\frac{d}{d\beta} \frac{4(\beta - 4)(\beta - 3)}{\beta^4} = \frac{-4(2\beta^2 - 21\beta + 48)}{\beta^5}$$

Then we have to set this to 0:

$$\frac{-4(2\beta^2 - 21\beta + 48)}{\beta^5} = 0$$

And find that there are two solutions:

$$\beta_1 = \frac{21}{4} - \frac{\sqrt{57}}{4}$$

$$\beta_2 = \frac{21}{4} + \frac{\sqrt{57}}{4}$$

Then, to find whether they are maxima or minima, we compute derivative again, and solve with $\beta = \beta_1$ and $\beta = \beta_2$. With this, we find that the $\underset{\beta}{argmax}$ is at:

$$\beta = \frac{21}{4} + \frac{\sqrt{57}}{4}$$

(b)

To perform this statistical test, we have to go through the six steps:

- A model:
The model is given as $X \sim \text{Bin}(n, \theta)$, with $n = 20$
- A hypothesis:
A hypothesis is given as well: $H_0 : \theta = 0.5, H_1 : \theta \neq 0.5$
- A test statistic and distribution:
- A significance level:
The significance level is also given as $\alpha = 0.05$
- A rejection region:
A rejection region is given as: $\mathcal{R} = \{0, \dots, a\} \cup \{20 - a, \dots, 20\}$
- Compute test statistic:

3. Exercise

(a)

This exercise has been completed in the attached *K_Means.ipynb* Jupyter Notebook. The most relevant bits of source code can be seen here as well:

```
data = np.loadtxt("old_faithful.csv", delimiter=',', skiprows=1)

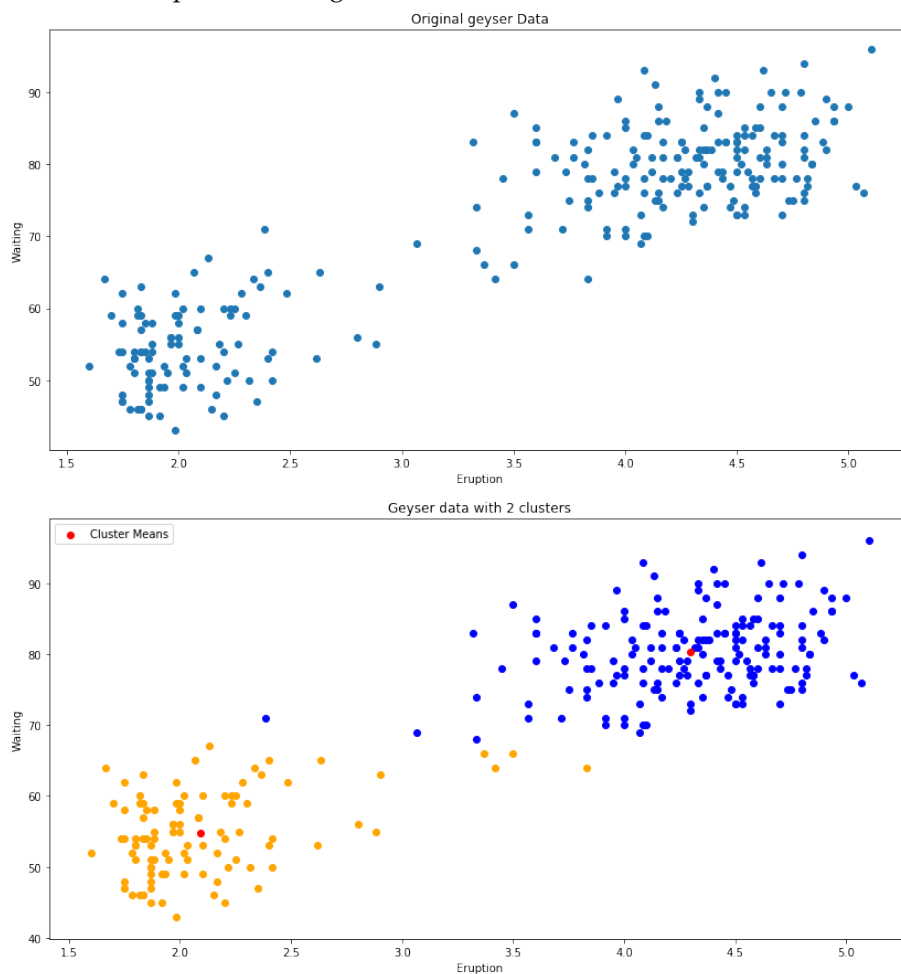
meanData = KMeans(n_clusters=2, max_iter=30, seed=0)
meanData.fit(data)

for i in range(data.shape[0]):
    if (meanData.cluster_assignments[i] == 0):
        plt.scatter(data[i,0],data[i,1], color='blue')
    else:
        plt.scatter(data[i,0],data[i,1], color='orange')
```

The 2 cluster means, calculated using the Jupyter Notebook, look as follows:

(4.29793023, 80.28488372)
(2.09433, 54.75)

And the resulting images, with, with each cluster in their own colour, and the cluster means plotted as singular red dots:



It is relevant to note that since the "Waiting" axis is significantly larger than the "Eruption" axis, it has a large effect on where the clusters are located. This can be seen especially on the leftmost blue points, and the rightmost orange points, which a human might expect to be the other colour.

(b)

This exercise has been completed in the attached *K_Means.ipynb* Jupyter Notebook. The most relevant bits of source code can be seen here as well, though all comments are left out. The final version is also different, as the function has to accommodate larger images. The function remains unchanged for this image though:

```
def makeNewImg(image, n_clusters, max_iter, seed):
    imgx = image.shape[0]
    imgy = image.shape[1]
    imgRGB = image.shape[2]
    model = KMeans(n_clusters, max_iter, seed)
    image = image.reshape((imgx*imgy),imgRGB)
    model.fit(image)

    retval = np.zeros(((imgx*imgy),imgRGB))

    for i in range(n_clusters):
        for j in range(imgRGB):
            model.cluster_means[i,j] = model.cluster_means[i,j]/255

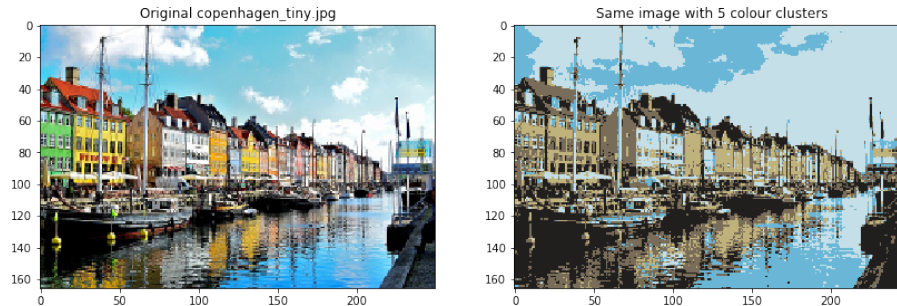
    for i in range(retval.shape[0]):
        retval[i] = model.cluster_means[model.cluster_assignments[i]]

    retval = retval.reshape((imgx,imgy,imgRGB))
    return retval, model.cluster_means
```

The 5 cluster means, calculated using the Jupyter Notebook, look as follows:

```
(0.41773344, 0.71691352, 0.84028094)
(0.13106363, 0.12181295, 0.11628755)
(0.47558986, 0.42848816, 0.35511877)
(0.76865056, 0.87663444, 0.90978927)
(0.75330927, 0.69978601, 0.48412478)
```

The assigned cluster for each pixel has been calculated, and the new image has been plotted:



(c)

This exercise has been completed in the attached *K_Means.ipynb* Jupyter Notebook. The most relevant bits of source code can be seen here as well, though all comments are left out. It should be noted that the code for this is the exact same as in the previous subexercise, except the "model.fit(image)" line has been exchanged for the following. The full source code is in the notebook:

```
if (imgx*imgy > 50000):
    modLen = 5000

    temp = np.arange(0, (imgx*imgy), 1).reshape((imgx*imgy))
    smallImgInd = np.random.choice(temp, modLen, False)
    smallImg = np.zeros((modLen, imgRGB))

    for i in range(modLen):
        smallImg[i] = image[smallImgInd[i]]

    model.fit(smallImg)
    model.cluster_assignments = model.assign_to_clusters(image, model.cluster_means)
else:
    model.fit(image)
```

The 16 cluster means, calculated using the Jupyter Notebook, look as follows. It should be noted that the cluster means have been calculated by choosing 5000 random pixels in the image, and fitting the model to those instead of the

entire image:

(0.04933315, 0.03762928, 0.03806449)
(0.57470588, 0.63714795, 0.57853832)
(0.23040784, 0.24485490, 0.24700392)
(0.55316917, 0.34053808, 0.13857729)
(0.43970764, 0.43151311, 0.37939419)
(0.99238280, 0.99199548, 0.98666989)
(0.94344192, 0.97583001, 0.98309507)
(0.41931736, 0.49164851, 0.55027838)
(0.82408510, 0.91116933, 0.93808050)
(0.73465750, 0.58706392, 0.38375350)
(0.71545709, 0.84472456, 0.88979713)
(0.55624748, 0.79929274, 0.88076591)
(0.33491397, 0.73982593, 0.90696279)
(0.73745915, 0.72773693, 0.68999183)
(0.93933934, 0.81688748, 0.27836071)
(0.91616099, 0.88028896, 0.78158927)

The assigned cluster for each pixel has been calculated, and the new image has been plotted:

